

★ Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



GENERATIVE AI SERIES

Ollama — Build a ChatBot with Langchain, Ollama & Deploy on Docker

Working with Ollama to run models locally, build LLM applications that can be deployed as docker containers.



A B Vijay Kumar  · [Follow](#)

5 min read · Feb 21, 2024



499



7



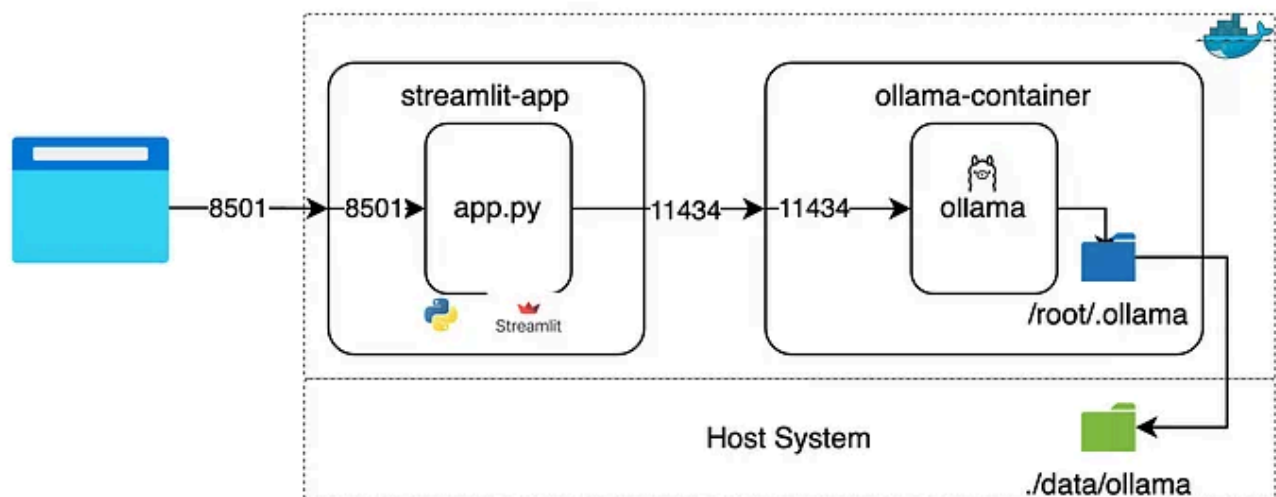
This blog is an ongoing series on GenerativeAI and is a continuation of the previous blogs. In this blog series, we will explore Ollama, and build applications that we can deploy in a distributed architectures using docker.

Ollama is a framework that makes it easy to run powerful language models on your own computer. Please refer to [Ollama — Brings runtime to serve LLMs everywhere.](#) | [by A B Vijay Kumar](#) | [Feb, 2024](#) | [Medium](#) for an introduction to Ollama. In this blog we will be building the langchain application and deploying on Docker.

Langchain Chatbot application for Ollama

Let's build the chatbot application using Langshan, to access our model from the Python application, we will be building a simple Steamlit chatbot application. We will be deploying this Python application in a container and will be using Ollama in a different container. We will build the infrastructure using docker-compose. If you do not know how to use docker, or docker-compose, please go through some tutorials on internet, before you go any further.

The following picture shows the architecture of how the containers interact, and what ports they will be accessing.



We will build 2 containers,

- Ollama container will be using the host volume to store and load the models (/root/.ollama is mapped to the local ./data/ollama). Ollama container will listen on 11434 (external port, which is internally mapped to 11434)
- Streamlit chatbot application will listen on 8501 (external port, which is internally mapped to 8501).

Before we start coding, lets setup a Python virtual environment.

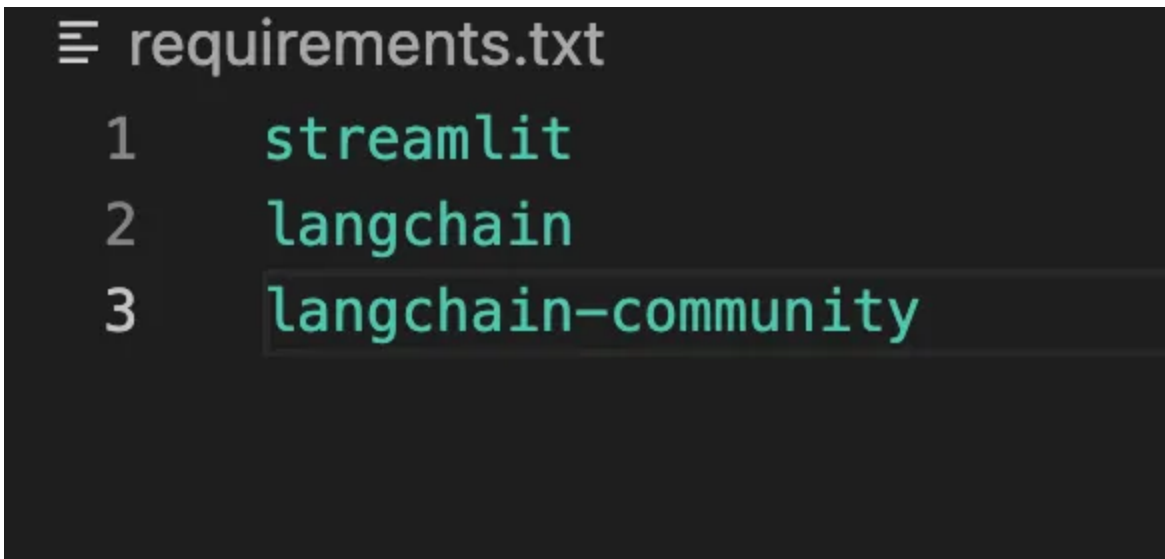
```
python3 -m venv ./ollama-langchain-venv
source ./ollama-langchain-venv/bin/activate
```

The following is the source code for streamlit application.

```
1  from langchain_community.llms import Ollama
2  import streamlit as st
3  from langchain.callbacks.streaming_stdout import StreamingStdOutCallbackHandler
4
5  llm = Ollama(model="phi:latest", base_url="http://ollama-container:11434", verbose=True)
6
7  def sendPrompt(prompt):
8      global llm
9      response = llm.invoke(prompt)
10     return response
11
12     st.title("Chat with Ollama")
13     if "messages" not in st.session_state.keys():
14         st.session_state.messages = [
15             {"role": "assistant", "content": "Ask me a question !"}
16         ]
17
18     if prompt := st.chat_input("Your question"):
19         st.session_state.messages.append({"role": "user", "content": prompt})
20
21     for message in st.session_state.messages:
22         with st.chat_message(message["role"]):
23             st.write(message["content"])
24
25     if st.session_state.messages[-1]["role"] != "assistant":
26         with st.chat_message("assistant"):
27             with st.spinner("Thinking..."):
28                 response = sendPrompt(prompt)
29                 print(response)
30                 st.write(response)
31                 message = {"role": "assistant", "content": response}
32                 st.session_state.messages.append(message)
```

This is very similar source code as I have built in my previous blogs. You can refer to my other blog [Retrieval Augmented Generation\(RAG\) — Chatbot for documents with LlamaIndex | by A B Vijay Kumar | Feb, 2024 | Medium](#) for details on how this code works. The main difference is we are using `ollama` and calling the model through Ollama Langchain library (which is part of `langchain_community`)

Let's define the dependencies in `requirements.txt`.



```
≡ requirements.txt
1  streamlit
2  langchain
3  langchain-community
```

Let's now define a Dockerfile to build the docker image of the Streamlit application.

```

1 FROM python:latest
2
3 # Create app directory
4 WORKDIR /app
5
6 # Copy the files
7 COPY requirements.txt ./
8 COPY app.py ./
9
10 #install the dependencies
11 RUN pip install --upgrade pip
12 RUN pip install -r requirements.txt
13
14 EXPOSE 8501
15 ENTRYPOINT ["streamlit", "run", "app.py", "--server.port=8501", "--server.address=0.0.0.0"]

```

We are using the python docker image, as the base image, and creating a working directory called `/app`. We are then copying our application files there, and running the `pip` installs to install all the dependencies. We are then exposing the port 8501 and starting the `streamlit` application.

We can build the docker image using `docker build` command, as shown below.

```

(=|docker-desktop:N/A)(base) → ollama-langchain docker build . -t abvijaykumar/ollama-langchain:0.1
[+] Building 2.1s (12/12) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring Dockerfile: 292B
=> [internal] load metadata for docker.io/library/python:latest
=> [auth] library/python:pull token for registry-1.docker.io
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [1/5] FROM docker.io/library/python:latest@sha256:e83d1f4d0c735c7a54fc9d0a3c2a6c58473a3b3de88fc67ba3d342ee73cfc89d
=> [internal] load build context
=> => transferring context: 63B
=> CACHED [2/5] WORKDIR /app
=> CACHED [3/5] COPY requirements.txt ./
=> CACHED [4/5] COPY app.py ./
=> CACHED [5/5] RUN pip install --upgrade pip
=> CACHED [6/5] RUN pip install -r requirements.txt
=> exporting to image
=> => exporting layers
=> => writing image sha256:d5978cad7e70a59978a57468f29f58de90d69c5079cfa95779c75cd71383d74
=> => naming to docker.io/abvijaykumar/ollama-langchain:0.1

```

View build details: [docker-desktop://dashboard/build/desktop-linux/desktop-linux/io220bm3t6z71vr34yaumu359](#)

You should be able to check if the Docker image is built, using `docker images`

[Open in app](#)

Medium

Search

Write



```
(*|docker-desktop:N/A)(base) → ollama-langchain docker images
REPOSITORY                                IMAGE ID      CREATED        SIZE        TAG
abvijaykumar/ollama-langchain             d5978cadb7e7  3 minutes ago  1.58GB      0.1
ollama/ollama                             c5761700ebce  28 hours ago   391MB      latest
```

Let's now build a docker-compose configuration file, to define the network of the Streamlit application and the Ollama container, so that they can interact with each other. We will also be defining the various port configurations, as shown in the picture above. For Ollama, we will also be mapping the volume, so that whatever models are pulled, are persisted.

```
docker-compose.yml - The Compose specification establishes a standard for the
1  version: '3'
2  services:
3    ollama-container:
4      image: ollama/ollama
5      volumes:
6        - ./data/ollama:/root/.ollama
7      ports:
8        - 11434:11434
9    streamlit-app:
10     image: abvijaykumar/ollama-langchain:0.2
11     ports:
12       - 8501:8501
13
```



We can bring up the applications by running the docker-compose up command, once you execute docker-compose up, you should be able to see that both the containers start running, as shown in the screenshot below.

```
([docker-desktop:N/A](base) → ollama-langchain docker-compose up
[+] Running 3/0
  ✓ Network ollama-langchain_default          Created
  ✓ Container ollama-langchain-streamlit-app-1 Created
  ✓ Container ollama-langchain-ollama-container-1 Created
Attaching to ollama-langchain-ollama-container-1, ollama-langchain-streamlit-app-1
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:01.193Z level=INFO source=images.go:706 msg="total blobs: 6"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:01.195Z level=INFO source=images.go:713 msg="total unused blobs removed: 0"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:01.196Z level=INFO source=routes.go:1014 msg="Listening on [::]:11434 (version 0.1.25)"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:01.197Z level=INFO source=payload_common.go:107 msg="Extracting dynamic libraries..."
ollama-langchain-streamlit-app-1 | Collecting usage statistics. To deactivate, set browser.gatherUsageStats to False.
ollama-langchain-streamlit-app-1 | You can now view your Streamlit app in your browser.
ollama-langchain-streamlit-app-1 | URL: http://0.0.0.0:8501
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.705Z level=INFO source=payload_common.go:146 msg="Dynamic LLM libraries [cpu cuda_v11]"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.705Z level=INFO source=gpu.go:94 msg="Detecting GPU type"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.705Z level=INFO source=gpu.go:262 msg="Searching for GPU management library libnvidia-ml.so"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.706Z level=INFO source=gpu.go:308 msg="Discovered GPU libraries: []"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.706Z level=INFO source=gpu.go:262 msg="Searching for GPU management library librocm-smi64.so"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.706Z level=INFO source=gpu.go:308 msg="Discovered GPU libraries: []"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.706Z level=INFO source=cpu_common.go:18 msg="CPU does not have vector extensions"
ollama-langchain-ollama-container-1 | time=2024-02-17T18:56:03.707Z level=INFO source=routes.go:1037 msg="no GPU detected"
```

you should be able to see the containers running by executing docker-compose ps command as shown below.

```
([docker-desktop:N/A](base) → ollama-langchain docker-compose ps
NAME                                IMAGE                                COMMAND                                SERVICE    CREATED        STATUS        PORTS
ollama-langchain-ollama-container-1 ollama/ollama                      "/bin/ollama serve"                  ollama-container    About a minute ago    Up About a minute    0.0.0.0:11434->11434/tcp
ollama-langchain-streamlit-app-1    abvijaykumar/ollama-langchain:0.2 "streamlit run app.py"                streamlit-app        About a minute ago    Up About a minute    0.0.0.0:8501->8501/tcp
```

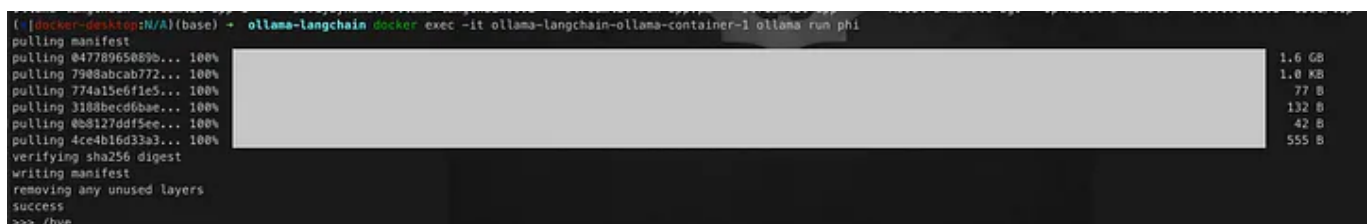
We should be able to check, if ollama is running by calling http://localhost:11434, as shown in the screenshot below.



Let's now download the required model, by logging into the docker container using the docker exec command as shown below.


```
docker exec -it ollama-langchain-ollama-container-1 ollama run phi
```

Since we are using the model phi, we are pulling that model and testing it by running it. you can see the screenshot below, where the phi model is downloaded and will start running (since we are using `-it` flag we should be able to interact and test with sample prompts)

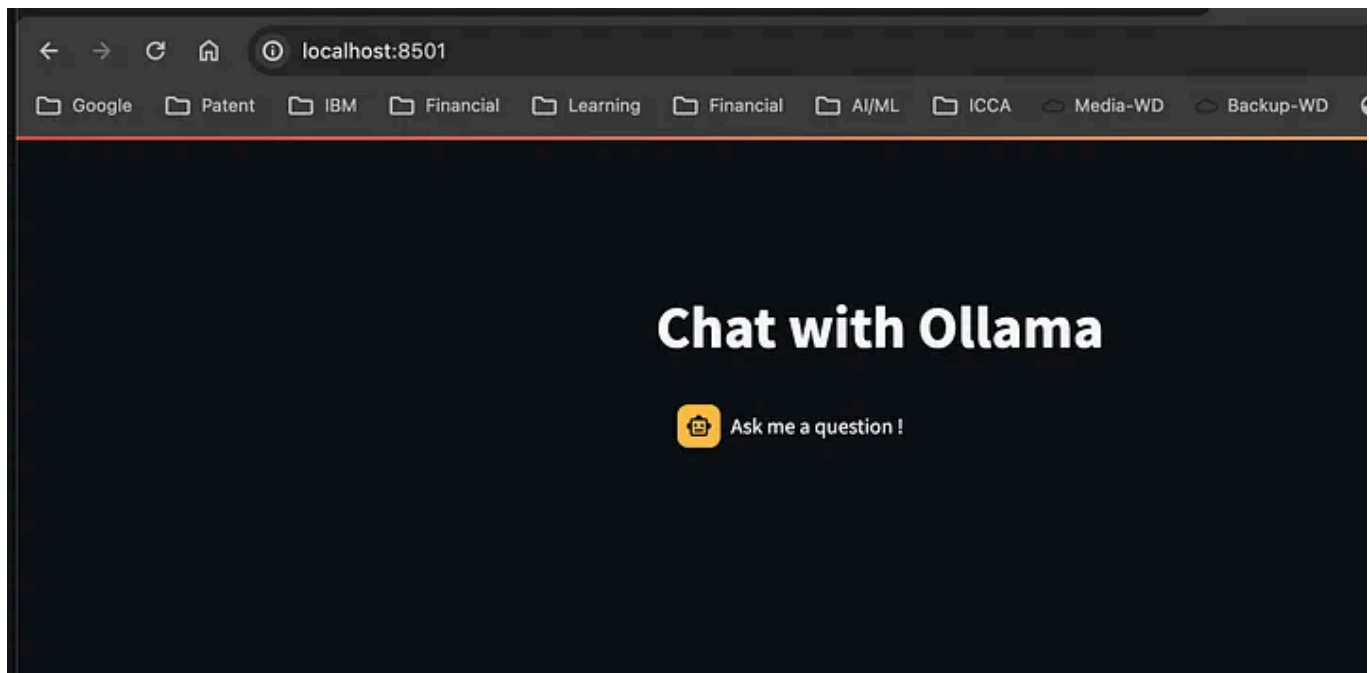


you should be able to see the downloaded model files and manifests in your local folder `./data/ollama` (which is internally mapped to `/root/.ollama` for

the container, which is where Ollama looks for the downloaded models to serve)



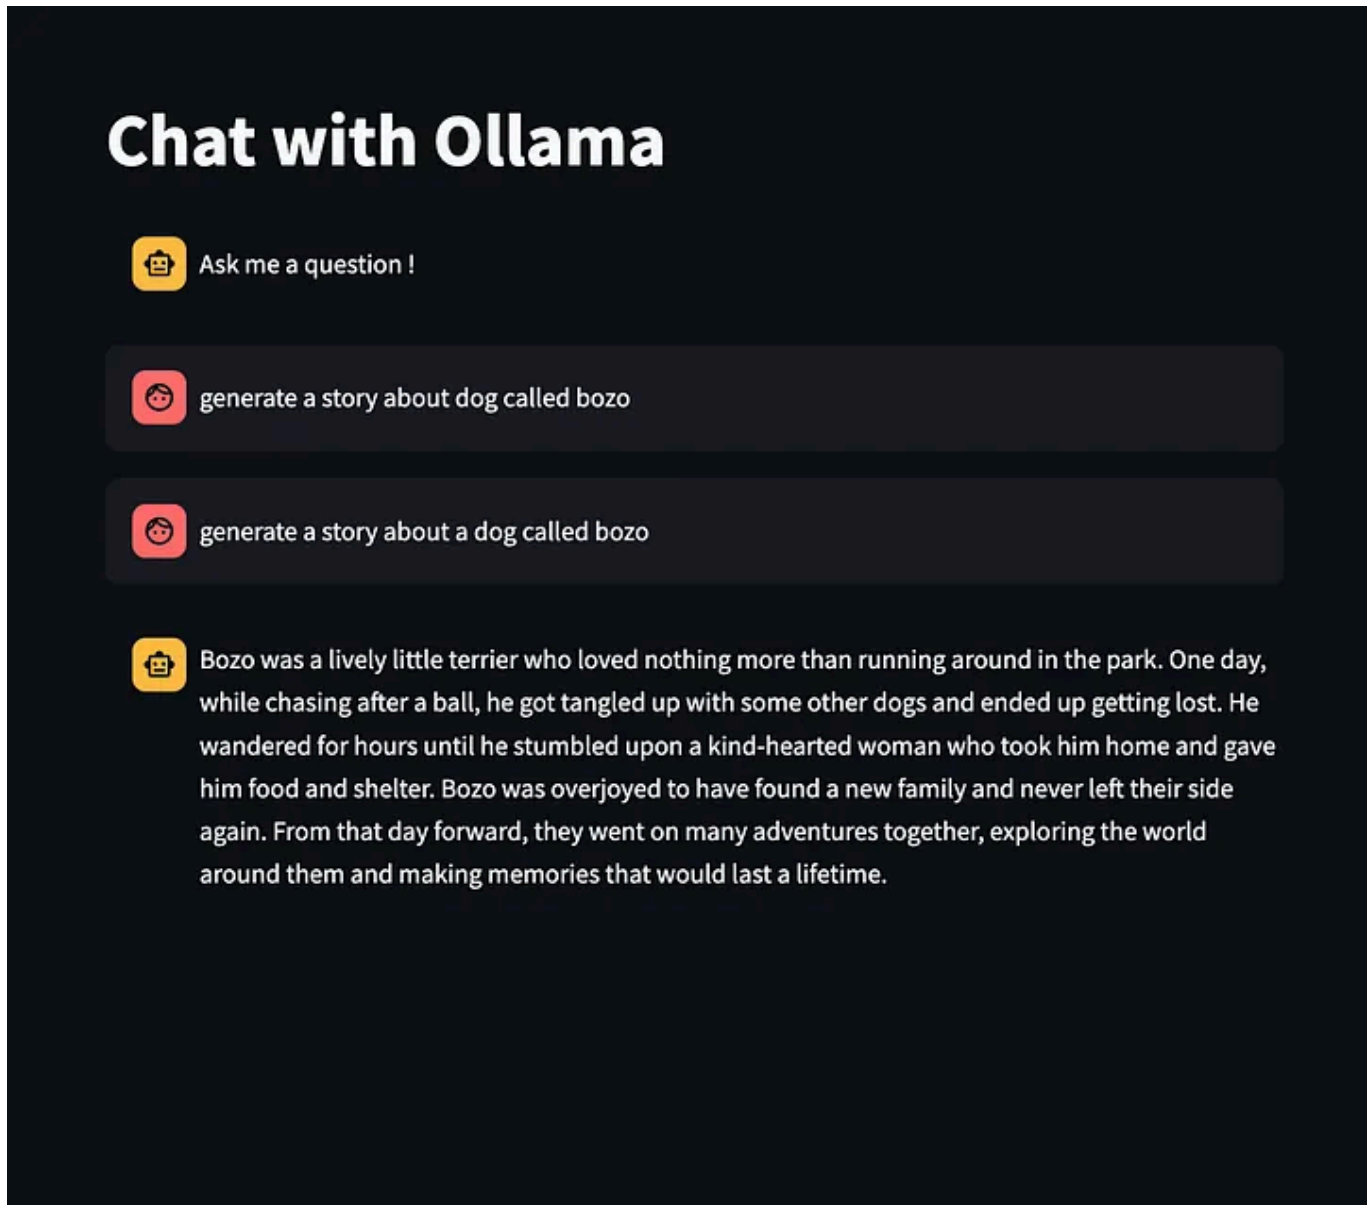
Lets now run access our streamlit application by opening <http://localhost:8501> on the browser. The following screenshot shows the interface



Lets try to run a prompt “generate a story about dog called bozo”. You shud be able to see the console logs reflecting the API calls, that are coming from our Streamlit application, as shown below

```
ollama-langchain-ollama-container-1 | llama_new_context_with_model: KV self size = 640.00 MiB, K (f16): 320.00 MiB, V (f16): 320.00 MiB
ollama-langchain-ollama-container-1 | llama_new_context_with_model: CPU input buffer size = 10.01 MiB
ollama-langchain-ollama-container-1 | llama_new_context_with_model: CPU compute buffer size = 163.00 MiB
ollama-langchain-ollama-container-1 | llama_new_context_with_model: graph splits (measure): 1
ollama-langchain-ollama-container-1 | time=2024-02-17T07:29:20.820Z level=INFO source=dyn_ext_server.go:156 msg="Starting llama main loop"
ollama-langchain-ollama-container-1 | [GIN] 2024/02/17 - 07:29:20 | 200 | 492.419875ms | 127.0.0.1 | POST | "/api/chat"
ollama-langchain-ollama-container-1 | [GIN] 2024/02/17 - 07:29:46 | 200 | 5.686822918s | 172.22.0.2 | POST | "/api/generate"
```

You can see in below screenshot, the response, I got for the prompt I sent



you can bring down the deployment by calling docker-compose down

The following screenshot shows the output

```
(+ | docker-desktop:N/A) (base) → ollama-langchain docker-compose down
[+] Running 3/0
✓ Container ollama-langchain-streamlit-app-1      Removed
✓ Container ollama-langchain-ollama-container-1    Removed
✓ Network ollama-langchain_default                 Removed
```

There you go. It was super fun, working on this blog getting Ollama to work with Langchain, and deploying them on Docker using Docker-Compose

Hope this was useful. I will be back with more experiments, in the meantime, have fun, and keep coding!!! see you soon!!!

you can access the full source code in my GitHub here. [abvijaykumar/ollama-langchain \(github.com\)](https://github.com/abvijaykumar/ollama-langchain).

References

- [Ollama](#)
- [Docker Compose overview](#) | [Docker Docs](#)
- [Docker Docs](#)
- [Ollama — Brings runtime to serve LLMs everywhere.](#) | [by A B Vijay Kumar](#) | [Feb, 2024](#) | [Medium](#)

Langchain

Generative Ai Tools

Ollama

Docker

Kubernetes

**Written by A B Vijay Kumar**

2.2K Followers · 209 Following

Follow



IBM Fellow, Master Inventor, Mobile, RPi & Cloud Architect & Full-Stack
Programmer

Responses (7)



Akhan

What are your thoughts?



Shruti Verma she/her

Mar 13



Thank you very much for sharing this guide; it helped me a lot! Give ServBay a try to effortlessly tackle all your development needs—it's fantastic!



1



1 reply

[Reply](#)



Scott Davies

Jul 6, 2024



Langshan

Langchain?



1

[Reply](#)



Ghaniasarwar

Feb 1



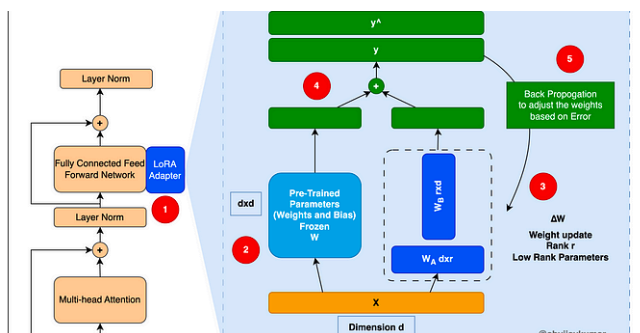
Can you please explain what is base url and how to get it i am using llama 3.1 model



1 reply

[Reply](#)[See all responses](#)

More from A B Vijay Kumar



A B Vijay Kumar

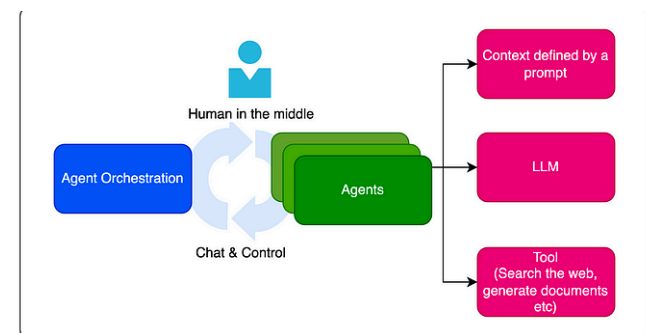
Fine Tuning LLM: Parameter Efficient Fine Tuning (PEFT)—...

Parameter Efficient Fine Tuning—LoRA, QLoRA—Concepts

Aug 9, 2023

608

9



A B Vijay Kumar

Multi-Agent Architectures

Multi-Agent systems are LLM applications that are changing the automation landscape...

Mar 24, 2024

631

6





A B Vijay Kumar



Multi-Agent System—Crew.AI

Multi-Agent systems are LLM applications that are changing the automation landscape...

Apr 17, 2024



462



7



A B Vijay Kumar



Fine Tuning LLM: Parameter Efficient Fine Tuning (PEFT)—...

Parameter Efficient Fine Tuning—LoRA, QLoRA—Hands-On

Sep 1, 2023



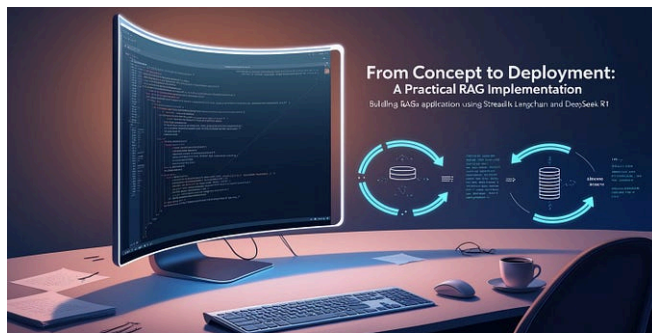
350



4

[See all from A B Vijay Kumar](#)

Recommended from Medium






 In Level Up Coding by Manasab Rao

Building a RAG Application using Streamlit Langchain and DeepSee...

From Concept to Deployment: A Practical RAG Implementation

★ Mar 3 🖱 1




 Anthony Sun

How to build chatbot with a local LLM in 5 minutes

Generated by ChatGPT

★ Oct 30, 2024 🖱 53




 In Data Science Collective by Gao Dalie (高達烈)

LangGraph + MCP + Ollama: The Key To Powerful Agentic AI

In this story, I have a super quick tutorial showing you how to create a multi-agent...


★ Mar 28 🖱 1.1K 💬 7





 Hi! How may I help you?

 Who are you?

 I am Hermione Jean Granger, one of the brightest witches of my age and a close friend of Harry Potter and Ron Weasley.

 Can you help me brew a polyjuice potion?


 I can guide you through the process, but I must warn you that brewing Polyjuice Potion is quite complex and requires precise measurements. To begin:

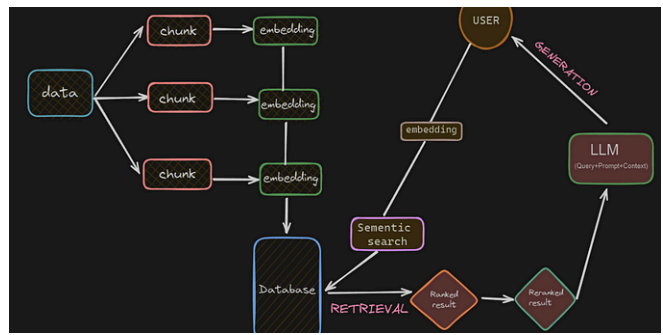
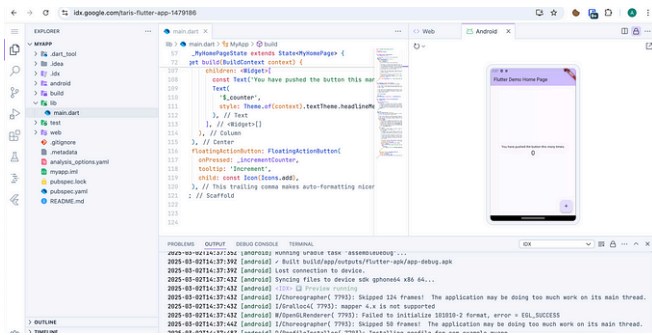
 Moirangthem Gelson Singh

Building a Local Chat Application with Streamlit, Ollama and Llama...

In this blog, we'll create a simple and fun chat application using Streamlit and Llama 3.2...

Dec 18, 2024 🖱 6





In Coding Beauty by Tari Ibaba

This new IDE from Google is an absolute game changer

This new IDE from Google is seriously revolutionary.



Mar 11



4.2K



231



Nov 30, 2024



93



See more recommendations



DhanushKumar

RAG with LLaMA Using Ollama: A Deep Dive into Retrieval...

The landscape of AI is evolving rapidly, and Retrieval-Augmented Generation (RAG)...